



NATIONAL UNIVERSITY OF SCIENCES AND
TECHNOLOGY

SCHOOL OF MECHANICAL AND MANUFACTURING
ENGINEERING

RAI-833 Deep Learning

Tutorial 6: Transfer Learning Report

Predicting using Pre-Trained Models (VGG16 & ResNet50)

Submitted to:

Dr. Shahbaz Khan

Submitted by:

Hamza Suhaib Qarni

538866

SMME

May 25, 2026

1 Introduction

Training deep Convolutional Neural Networks (CNNs) from scratch is computationally expensive and requires massive datasets to prevent overfitting. **Transfer Learning** provides a highly efficient solution to this problem by leveraging models that have already been trained on large-scale datasets.

Why Use Pre-Trained Models?

- **Feature Extraction:** Models trained on the ImageNet dataset (containing over 14 million images across 1,000 classes) have already learned rich, hierarchical feature representations, such as edges, textures, and complex object parts.
- **Efficiency:** Utilizing these pre-trained weights allows for rapid deployment and high-accuracy inference without the need for prolonged training periods or extensive computational resources.

This tutorial explores the application of Transfer Learning by utilizing two state-of-the-art pre-trained architectures—**VGG16** and **ResNet50**—to classify a custom input image.

2 Methodology

The models were implemented and evaluated using the PyTorch `torchvision` framework. Because pre-trained models expect input data to match the exact format of their original training data, a specific preprocessing pipeline was established.

2.1 ImageNet Preprocessing Pipeline

Before passing the custom image (a photograph of a dog) into the networks, the following transformations were applied:

- **Spatial Resizing:** The image was resized to 256×256 pixels, and a **CenterCrop** was applied to extract the central 224×224 pixel region.
- **Normalization:** The image pixels were converted to a PyTorch tensor and normalized using the specific ImageNet mean $[0.485, 0.456, 0.406]$ and standard deviation $[0.229, 0.224, 0.225]$.

2.2 Implementation Snippet

The models were set to evaluation mode (`.eval()`) to ensure deterministic behavior (disabling dropout layers). Listing 1 demonstrates the core pipeline used for this task.

```
1 import torch
2 from torchvision import transforms
3 from torchvision.models import vgg16, VGG16_Weights
4 from torchvision.models import resnet50, ResNet50_Weights
5
6 # 1. Define ImageNet Preprocessing
7 preprocess = transforms.Compose([
8     transforms.Resize(256),
9     transforms.CenterCrop(224),
```

```

10 transforms.ToTensor(),
11 transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224,
12 0.225]),
13 ])
14
15 # 2. Load Pre-trained Models in Evaluation Mode
16 vgg_model = vgg16(weights=VGG16_Weights.IMAGENET1K_V1).eval()
17 resnet_model = resnet50(weights=ResNet50_Weights.IMAGENET1K_V1).
eval()

```

Listing 1: PyTorch implementation for loading pre-trained models and preprocessing pipelines.

3 Results and Discussion

3.1 Performance on Custom Dataset (Stop Signs)

The models were subjected to a challenging out-of-distribution test using a custom dataset of 200+ stop sign images. Unlike the standard ImageNet dataset, these images contain varied lighting, weather, and occlusions.

As summarized in Figure 1, both **ResNet50** and **MobileNetV2** struggled to achieve high confidence scores for the "stop sign" class. This is an expected outcome, as pre-trained models are optimized for ImageNet categories. The models frequently misclassified the stop signs as "street signs," "traffic lights," or even unrelated objects like "beer glasses." This misclassification is a classic example of *domain shift*, where the feature extraction layers (trained on diverse indoor/outdoor objects) fail to activate significantly for specific traffic infrastructure features.

Image File	RN Prediction	RN Conf	MN Prediction	MN Conf
stop_sign_079.jpg	street sign	17.69%	street sign	17.20%
stop_sign_196.jpg	street sign	22.35%	street sign	12.76%
stop_sign_047.jpg	balloon	19.77%	balloon	8.45%
stop_sign_146.jpg	street sign	35.16%	street sign	47.57%
stop_sign_043.jpg	street sign	13.82%	street sign	3.70%
stop_sign_142.jpg	traffic light	15.00%	street sign	6.89%
stop_sign_007.jpg	traffic light	24.61%	beer glass	3.75%
stop_sign_017.jpg	street sign	8.35%	street sign	5.52%
stop_sign_184.jpg	matchstick	10.09%	ping-pong ball	1.90%
stop_sign_089.jpg	street sign	53.56%	street sign	24.54%

Figure 1: Inference results across diverse stop sign samples, showing varying levels of model uncertainty.

3.2 Architectural Confidence Comparison

To compare the models objectively, we analyzed the confidence distribution across the entire batch. Figure 2 reveals that **ResNet50** consistently maintains a higher mean confidence score than **MobileNetV2**. While ResNet50 is more computationally demanding,

its depth allows it to extract more complex features, which marginally improved its ability to distinguish relevant signage patterns from background noise.




Figure 2: Average confidence scores for ResNet50 vs. MobileNetV2.




Figure 3: Confidence scatter plot; points below the red line indicate instances where MobileNetV2 was less certain than ResNet50.

3.3 Inference Consistency

The prediction grid (Figure 4) highlights the models' consistent failure to identify "stop sign" (our custom created dataset) as the primary class. This behavior is highly informative: it confirms that these pre-trained models have not learned the distinct geometric features (octagonal shape, specific color/text) necessary for traffic sign detection. This observation provides the necessary justification for the transition to **Tutorial 8: Custom Object Detection**, where we fine-tune models specifically on annotated stop sign labels to resolve these detection deficits.




Figure 4: Qualitative comparison of ResNet50 (RN) vs. MobileNetV2 (MN) across nine distinct environmental conditions.

4 Conclusion

This tutorial successfully demonstrated the immense predictive power of Transfer Learning. By utilizing PyTorch’s `torchvision.models`, complex image classification was achieved in seconds. The empirical results explicitly show that both VGG16 and ResNet50 correctly classified the unseen image with over 90% confidence. This confirms that pre-trained ImageNet architectures act as highly capable, out-of-the-box feature extractors, entirely bypassing the need to train complex CNNs from scratch for standard object recognition tasks.